Amie L. Thomasson

"Philosophy as Conceptual Engineering " by Amie L. Thomasson (Keywords: Metaphysics; Metaphilosophy)

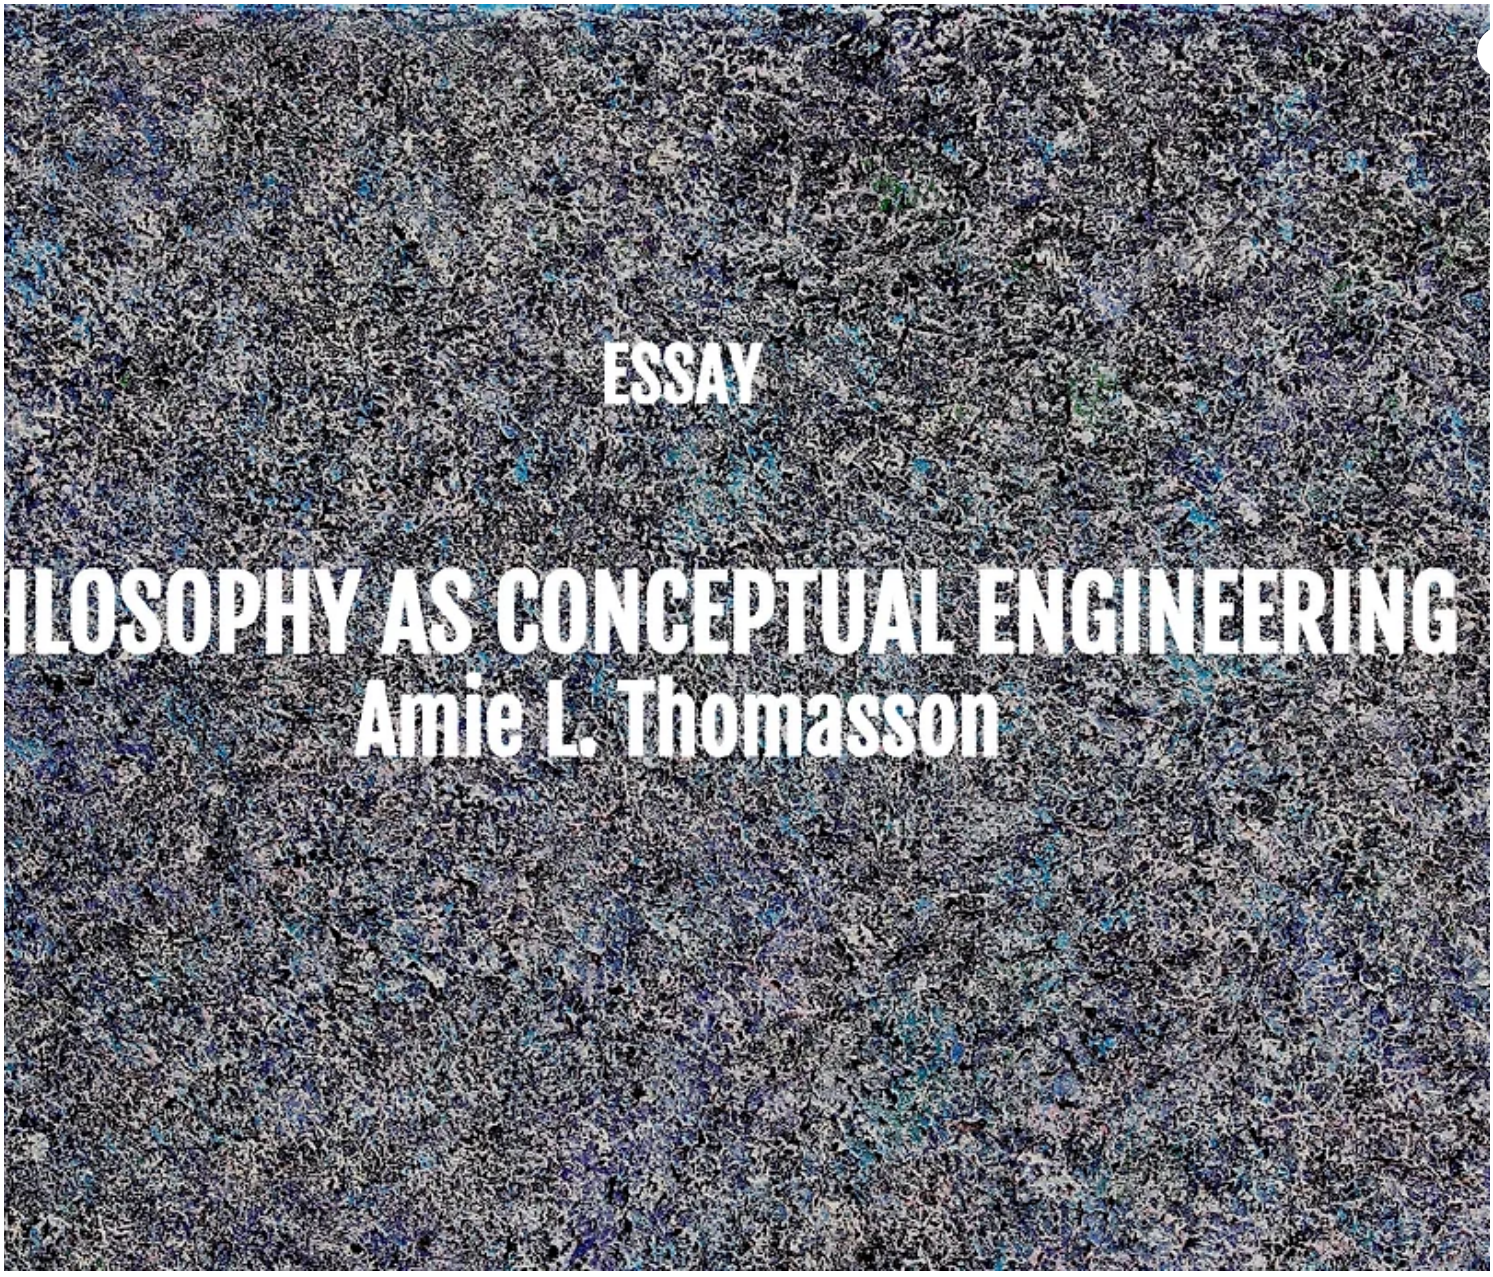

From *The Philosopher,* vol. 109, no. 3 ("Concept and Reality").

Philosophers tend to speak as if they are investigating and making discoveries about the world. Socrates asked, “What is justice?”, “What is courage?”; Aristotle described what he thought of as the “good life” for humans; Descartes claimed to discover the essence of minds, bodies, and himself; and Locke defended views about what property is, and what it takes to be the same person over time.

Many philosophers today – especially those in metaphysics – present themselves as making discoveries about what “really exists” or about the “fundamental features of the world”. If they think about concepts at all, they think of them only as ways we have of investigating reality, or communicating the results of our investigation.

But there is also a great deal of scepticism about the idea that philosophy can make discoveries about what exists, or about what the world is like. Stephen Hawking expressed this scepticism, writing (with Leonard Mlodinow in *The Grand Design*):

> ...people have always asked a multitude of questions: How can we understand the world in which we find ourselves? How does the universe behave? What is the nature of reality? Where did all this come from?.... Traditionally these are questions for philosophy, but philosophy is dead. Philosophy has not kept up with modern developments in science, particularly physics. Scientists have become the bearers of the torch of discovery in our quest for knowledge.

This is a familiar challenge for philosophy, and has been around ever since the natural sciences began to separate themselves from what was once called “natural philosophy”. As physics, chemistry, biology and other sciences became explicit about their own goals, methods, and criteria for success, the question remained: what (if anything) is left for philosophy?

> If we think of philosophy as aiming to tell us “how the world is”, it seems to be caught in a rivalry with the natural sciences – a rivalry it seems bound to lose.

As the natural sciences have gradually come to converge in their results and provide the basis for impressive technological and medical achievements, their credentials for claiming to tell us about the world have gotten pretty good. By contrast, the “theories” of philosophers on any given issue, far from converging, have only proliferated as time goes on, with ever new proposals and no agreement about what might resolve the standing debates. If we think of philosophy as aiming to tell us “how the world is”, it seems to be caught in a rivalry with the natural sciences – a rivalry it seems bound to lose.

In fact, the proliferation of philosophical views and failures of convergence alone might lead to scepticism about the ability of philosophy to “discover” truths about the world. Here’s a way of quantifying it. The Philosophical Survey, conducted in 2009, was a poll of 3,226 philosophers – the vast majority of whom were professors, PhDs, or grad students – asking their opinions on 30 core philosophical questions (questions about whether we have free will, about the existence of God, the objectivity of aesthetic value, when we have knowledge, and so on). Bryan Frances (himself a philosopher) did a statistical analysis of the results. Based on the diversity of answers given (and assuming we think there are correct answers to these philosophical questions), Frances argued that where definite answers were given, the average philosopher can be thought to get 47%-67% right. Hardly an impressive score! And these should be the experts. (I’m sure we all hope to get

more accuracy from our medical experts). As a result, he argued that we should suspend judgment indefinitely, refraining from believing anything on these issues. This sort of radical disagreement in philosophy might then lead us to a sceptical torpor – thinking that, undiscoverable as such facts are, we may as well give up trying and turn our attention elsewhere.

***

There is, of course, an alternative to thinking of philosophy as making worldly discoveries that would rival those of the sciences. That is to think that philosophy is not concerned with quasi-scientific discoveries, but rather with analyzing our language or concepts. Two of the great philosophical movements of the early to mid-twentieth century – Phenomenology and Ordinary Language Philosophy – took something like this approach, while philosophers as diverse as Husserl, Ayer, Wittgenstein, Ryle, Strawson and Austin took the work of philosophy to lie in analyzing our language, meaning, or conceptual scheme.

But doubts about meanings and analyticity, and a general frustration with this "limited" conception of what philosophy can do led to a great backlash against it. In his book *Four-Dimensionalism*, Ted Sider asked, "Who would prefer exploring our perhaps parochial conceptual scheme to exploring the fundamental features of reality?" while recent metaphysicians tend to characterize inquiries into our language or conceptual scheme as merely superficial, trivial, uninteresting, and unimportant. As a result, it is still common for contemporary philosophers to insist that they are not merely interested in our thought or talk; rather, they are interested in *the world.*

So, we seem left with two options:

1) Thinking of philosophy as engaged in investigating "facts" of the world leads to scepticism and to the threat of a rivalry with the natural sciences – a rivalry it seems bound to lose,
2) Thinking of philosophy as involved in investigating our language or concepts risks making it trivial, superficial, and uninteresting – and perhaps subject to its own rivalry with the investigations of linguistics or cognitive science.

***

But this view of the options leaves out something important. Philosophy hasn't been, and needn't be, concerned just with how we *do* think or talk. We can also ask how we *should* think and talk. Nor are philosophers constrained to just *use* our existing concepts to investigate the world. Instead, we can work to evaluate and reconstruct the concepts and language we use, and sometimes even to construct new concepts and terms to help us.

> Philosophy hasn't been, and needn't be, concerned just with how we *do* think or talk. We can also ask how we *should* think and talk.

Rudolf Carnap suggested an idea along these lines in 1950. Rather than thinking of philosophy as a rival to the sciences, he suggested, we can think of it as useful in clarifying, modifying and developing concepts for use in

the empirical work of the sciences. That is, philosophers may aim to *explicate* the concepts we need to do science: concepts such as “confirmation”, “probability”, “temperature” and “species”. Explication, on Carnap’s view, isn’t a matter of reporting how our concepts *do* work, but rather *improving them* by making them more precise and more fruitful for serving in generalizations and laws, and thus more useful for the sciences. With this, Carnap recognized the possibility of a project of conceptual engineering.

Recent work in philosophy has seen an explosion of interest in conceptual engineering, and a broadening of its project. In part, the recent interest in conceptual engineering comes from recognizing that it is not only our *scientific* concepts that might be improved. Sally Haslanger raised the question of whether social concepts, such as race and gender concepts, might be improved – not to serve our scientific theories, but rather to serve the cause of social justice. Conceptual engineering projects have broadened greatly in scope since Carnap’s beginnings, including not only attempts to (re-)engineer scientific concepts, but also social concepts, logical concepts, and other philosophical concepts.

But what is conceptual engineering? Think of conceptual engineering on analogy to civil or mechanical engineering. Engineers may be involved in constructing new bridges or roadways, or in designing new kinds of machinery or structures to serve new purposes. But it is also engineers we look to when we need to assess our old structures – whether a building remains stable after an earthquake, what repairs or maintenance an aircraft needs to safely carry passengers, or what reinforcements a roof needs to withstand a possible storm.

Similarly, conceptual engineering may focus on re-evaluating our extant concepts. We may need to examine our concept of truth, or freedom, or person. Do these concepts lead us into paradoxes, unresolved questions, or unacceptable conclusions? Or it may focus on reconstructing our old concepts in a changed context – in the current social and technological context, what concepts do we need of intelligence, privacy, information, disease, and so on? Or it may focus on constructing new concepts to cope with and often reshape our social reality or our methods of inquiry – introducing concepts such as sexual harassment, genocidal rape, gene, or autism.

If we think of philosophy as engaged in conceptual engineering, it no longer faces a rivalry with the sciences. As Carnap saw, it may work hand-in-hand with the sciences – with physics in asking how we should think of space and time, or with biology in asking what species concept we should use, or more broadly in asking how we should understand confirmation, evidence, and inference. Or it may work on concepts that are not the province of science at all, but still play central roles in human life – concepts such as art, freedom, responsibility, consciousness, or person.

Thinking of philosophy as engaged in conceptual engineering leads to a different way of thinking about the relation between concepts and reality, as we will no longer see concepts as simply transparent windows through which we may see reality, or report what we see. Instead, we can come to acknowledge the ways in which concepts can serve as tools to shape how we think and live.

***

How can changes in our conceptual scheme lead to changes in the world?

It takes only a little reflection to note the ways in which the great social movements of the last century have been tied up with revolutions in the language and concepts we use.

Our social concepts are typically tied not simply to “picking out” some group in the world, but to imposing certain norms for how members of that group are to be treated and regarded. Pejoratives (say for race or national origin) do not simply pick out a group of people, but derogate members of that group. Even terms that are not pejoratives may impose norms of treatments – think of why feminists have long emphasized the need to stop calling grown women “girls”.

> Our social concepts are typically tied not simply to “picking out” some group in the world, but to imposing certain norms for how members of that group are to be treated and regarded.

Sometimes terms and concepts enable the formation of whole systems of laws – consider the old “racial” terms such as “quadroon” and “octoroon” – used in imposing fine-grained discriminatory legal practices. Drop the concepts and it is no longer possible to develop or enforce such systems of laws and the fine-grained racial hierarchies they supported.

Sometimes changes of concepts go hand-in-hand with changes in laws, or in how laws are applied. As the criteria for marriage were reformed in many places, the old inference rules that gave entitlement to infer that one man and one woman were involved in any marriage were dropped as well. The concept was changed, and with it our social reality changed. For while the inferential connections to the gender of the married people changed, the connections to thousands of legal rights and entitlements tied to marriage remained, as did the more informal connections to gaining public social recognition of a life-shaping relationship. Similarly, the ways we formulate our concepts of disease and of disability also have immediate legal consequences – for what sorts of accommodation are required in schools, and what sorts of treatment are paid for by insurance.

Such changes aren’t limited to changing old terminology. Think also of the role conceptual and linguistic *innovation* has played in many of the social changes of recent decades. Introducing words (and concepts) like “sexual harassment” (coined by activists at Cornell University in 1975) and “hostile environment” was the key to getting certain patterns of behaviour identified, acknowledged as problematic, and as suitable for prohibition and prosecution. Introducing the term and concept of “genocidal rape” was crucial to the developing international laws for its prosecution (philosophers including Asja Armanda, Natalie Nenadic and Catherine MacKinnon were central in working to bring about this change).

Nor are the changes wrought by our concepts limited to legal or social changes. Changes in our concepts also go along with changes in our practices of investigation and treatment. Think of the changes in our terminology for emotional, behavioural, and cognitive differences, as we moved using from terms like “madman” or “lunatic” to speaking of people with “mental illness”, and from using terms like “idiot” or “moron” (both of which once had “technical” meanings) to talking instead of people with “developmental delays”, “special needs” or “cognitive disabilities”. Think of the difference between whether our psychological terms reflect a supernatural, medical, or social model of cognitive and behavioural differences – and of the differences these lead to in whether we engage in an attempted religious, biochemical, behavioural, or contextual treatment of these differences.

As Thomas Kuhn emphasized, conceptual change has also been central to scientific revolutions. Think of the conceptual changes in the meaning of "gene" that were crucial to the development of biology. Moreover, the ways we understand what should count as the same, or a different, species, makes an enormous difference to how we investigate and report on biodiversity, how we identify and count endangered species, and how we make plans for environmental protection or remediation. In physics, changes in how we understand space, time, and their relationship were essential to developing the theory of relativity.

All this change – in our social world, in legal rights and obligations, in investigative practices – comes about with changes in our concepts. And that is because our terms and concepts aren't just ways of pointing to individuals or transparently discovering how the world is. Our terms come with *norms* – often implicit rules for what we can infer from them. And these may be inferences about how individuals of different groups should be treated, what rights and obligations they have, or how conditions of various kinds are to be investigated and treated. Change the concepts, and you can change the inferences, and in turn you can change the social, legal, and even scientific practices.

> Change the concepts, and you can change the inferences, and in turn you can change the social, legal, and even scientific practices.

To allow that changes in our conceptual scheme can have worldly importance in this way is not to embrace a kind of idealist philosophical view that our minds "create reality". It is merely to acknowledge the ways in which our concepts are tied to all kinds of *norms* for how we think and live – and to acknowledge the enormous difference this may make for how we treat each other, how we run our social institutions, and how we conduct our scientific investigations.

***

A focus on conceptual engineering makes it clear that our concepts aren't merely transparent tools we use to make discoveries in the world. They are crucial to *shaping* the world, by shaping how we think and what we do.

We have already seen some of the ways in which re-evaluating concepts such as race, disability, disease, and privacy, have been socially important, as well as the impact of re-evaluating scientific concepts such as space, time, and species. All of these have been and remain important philosophical topics.

Historically, however, philosophy has often focused on even more *general* concepts that have long been central to human life: concepts like *person, freedom, right, art, knowledge,* and *justice.* Consider to what a great extent our lives are shaped by what sort of action counts as *free*; by what sort of political *freedom* a state should guarantee; by what sorts of *rights* should be enforced; by who counts as having *knowledge*; by what counts as *justice*; or by what qualifies as a work of *art* worthy of exhibition, study, and veneration.

Understanding, assessing, and re-engineering these concepts has always played a central role in philosophical work; and shaping our understanding and use of these concepts has had and can have an enormous impact on human life.

But wait, some might say, the classic philosophers don't present what they're doing as just aiming to understand and shape our *concepts* – they talk about the *world*. Plato talks about *what justice is,* John Locke about what *makes a person the same over time,* David Hume about *what liberty is,* Clive Bell and Arthur Danto about *what art is,* and so on.

It is certainly true that many philosophers simply *use* the relevant words in their discussions of what justice, persons, liberty, or art are. But often, even in our daily lives – in our political conversations, arguments with our friends, pub discussions – we use words not as a way of expressing disagreements about empirical facts of the world, but rather as ways of pressing for views about how we *should* use our terms.

Think of the debates that inevitably ensue with the Olympics, about whether Air Pistol or Synchronized Swimming are really sports. Or think of more serious debates about whether the Oklahoma City bombing was terrorism, or whether waterboarding is torture.

These are cases of what David Plunkett and Tim Sundell call "Metalinguistic Negotiation": where we *use* words implicitly in order to negotiate for how (or whether) these terms *ought to be* used. Such debates needn't arise from different beliefs about other facts (say about the skills or training of Air Pistol competitors), and there's no sense that further "discoveries" would resolve the debate. Nor are such debates resolved just by seeing how the relevant words ("sport", for example) are *actually* used or by acknowledging that the disputants are using a key word in different ways. For at bottom these disputes are not about what the *world is like* or about how words *are used*, but rather about how they *should be used.*

When John Locke argues that preserving a continuity of consciousness (not just of body) is essential to the same person continuing to exist over time, he knows that this doesn't match the standard practices of the time. But he argues that this is how we *should* come to think of persons, since person is "a forensic term, appropriating actions and their merit", and it is only right and just to reward and punish one who shares a continuous consciousness with the person who committed the act.

When Paul Edwards argues that there is no free will, he is not concerned with the fact that we typically call an action "free" if someone is doing what they want to do, without external constraint. He does so by arguing that taking people to be *free* means holding them *responsible*. But, he argues, we shouldn't treat anyone as responsible unless they originally *chose* their own character, and he argues that this condition is never met.

What such philosophers are doing can be seen as pushing for views about how we *should* think of persons or of freedom. And such debates aren't just "about words". They matter for who we blame, punish, give property and other rights to, for how we raise our children and treat our friends and our debtors.

***

Thinking of much traditional work of philosophy as work in conceptual engineering enables us to avoid the threat of a rivalry with science while also denying that treating philosophical work as conceptual work makes it shallow or unimportant.

But it also leaves us with two important matters to be addressed:

1) The Assessment Problem: How *can* we work out what concepts we should use, or how we should use them?

2) The Implementation Problem: Even if we figure that out, how could we make these changes in our conceptual scheme or language?

One way to make progress on the first question is to follow Wittgenstein in thinking of our language as like a toolbox. Our terms and concepts serve as tools – not just to report on or investigate the world, but to *do* many things: to greet, express our attitudes, communicate rules of games or rules of inference, make calculations, organize our laws and social lives, and much more. Once we think of our language and concepts in this way, we can see how to approach the assessment problem. As Matthieu Queloz has suggested, a good way to begin with terms of our current language is through *reverse engineering.* That is, we should begin by assessing what terms such as number terms or moral terms *do* that has made them a useful part of our conceptual repertoire (or at least useful for some!). Once we have a good idea of what the relevant *functions* are, we can do better at assessing whether these are functions we should retain or reject. Are they (like number terms) functioning in crucial ways to our keeping track of information, reasoning, and making inferences in science and elsewhere? Or are they (like pejoratives) functioning to derogate? We can also better assess whether the relevant terms and concepts are still serving their functions well (even in changed technological and social circumstances), or if we might do better with changing our concept of truth, or freedom, or death, or marriage, or intelligence.

Yet even if we can work out what we *should* do with our concepts, some have expressed scepticism about whether we *could* implement changes to our language or conceptual scheme. Herman Cappelen raises this as the "implementation problem" for conceptual engineering. But this, too, is not as difficult a task as critics have made out. If we think of our terms and concepts as tools, it is natural to also think of them as having *rules of use,* just as there are proper or improper ways to use a hammer, a fork, or a canoe. In the case of words, there are social rules for what inferences we may and should make, from the claim that a certain statement is *true,* that a law is *just,* or that an action was *free.* The problem of how we can change our language, then, can be seen as part of the general problem of how we can change *social norms* – those (often informal and unstated) rules for how we are to behave in a variety of circumstances. And while that is not easy, thanks in part to the work of philosophers like Christina Bicchieri we are now learning more and more about how that *can* be done through public argument or debate, formal stipulations by governmental or professional bodies, or the exemplary behaviour of trendsetters.

Of course, much work remains to be done in working out how we can best assess our language and concepts – how we can determine their functions and how well they are fulfilling them. Other work remains to be done in figuring out how to implement any needed changes, to help the desired conceptual changes "catch on". Work in linguistics may help with both of these problems – systemic functional linguistics may help us identify linguistic functions; historical linguistics may help us identify what processes underlie natural linguistic change, so that we can hope to exploit these in inducing change.

***

Conceptual engineering provides a promising and interesting route for reconceptualising central areas of philosophy – in a way that both avoids the threat of a rivalry with science and also avoids the feeling that philosophical work is trivial or uninteresting. (Other parts of philosophy, such as ethics, may not need any such reconceptualisation to avoid these problems.) This conceptual work can both make sense of much of what philosophy *has done,* and of what philosophy *can do,* as well as why it matters. In engaging in conceptual engineering, it is also crucial to make the *reasons* for these conceptual and linguistic choices explicit and open for examination. And that will require bringing upfront the reasons for these conceptual choices, rather

than pretending we are like scientists and making our own “discoveries”. It will require being transparent about what we are doing and abandoning the quasi-scientific “discovery” model of philosophy.

> We must go on using some conceptual system, so the worst thing we could do would be to give up thinking seriously about which concepts we should use, and how we should use them.

Thinking of philosophy as conceptual engineering rather than worldly discovery also enables us to avoid the feeling that the answers to philosophical questions are unknowable, so we should just give up trying. We must go on using some conceptual system, so the worst thing we could do would be to give up thinking seriously about which concepts we should use, and how we should use them. For how we think of, and speak about, our world really matters.

Amie L. Thomasson is the Daniel P. Stone Professor of Intellectual and Moral Philosophy at Dartmouth College. She is the author of more than 75 articles and of four books: Norms and Necessity *(2020),* Ontology Made Easy *(2015 – winner of the Sanders Book Prize),* Ordinary Objects *(2007), and* Fiction and Metaphysics *(1999).*

From *The Philosopher,* vol. 109, no. 3 ("Concept and Reality").