Open Research Online

Conceptual Engineering, Metasemantic Externalism and Speaker-Meaning

Journal Item

For guidance on citations see FAQs.

oro.open.ac.uk

Conceptual Engineering, Metasemantic Externalism and Speaker-Meaning

MARK PINDER

The Open University

Mark.Pinder@open.ac.uk

Abstract. What is the relationship between conceptual engineering and metasemantic externalism? Sally Haslanger has argued that metasemantic externalism justifies the seemingly counterintuitive consequences of her proposed conceptual revisions. But according to Herman Cappelen, metasemantic externalism makes conceptual engineering effectively impossible in practice. After raising objections to Haslanger's and Cappelen's views, I argue for a very different picture, on which metasemantic externalism bears very little on conceptual engineering. I argue that, while metasemantic externalism principally operates at the level of semantic-meaning, we should understand conceptual engineering to largely operate at the level of speaker-meaning. This 'Speaker-Meaning Picture' has two key benefits. Firstly, it makes conceptual engineering often easy in practice. Secondly, it suggests a new, normative response to the well-known objection that conceptual engineering serves only to change the subject.

1. Introduction

Many concepts, it seems, are deficient. They may be insufficiently precise for theoretical inquiry, they may obscure important distinctions, they may be offensive or politically contentious, or something else. To deal with this, one may seek to refine one's concepts. This is called *conceptual engineering*.

Conceptual engineering appears to have an ambivalent relationship with metasemantic externalism, the view that external factors—the environment, experts, historical chains of use, etc.—play a key role in determining the meanings of linguistic expressions. On the one hand, Sally Haslanger has argued that metasemantic externalism can justify the seemingly counterintuitive results of conceptual engineering. For example, given the revisions of MAN and WOMAN proposed in

Haslanger 2000, there ought not to be any men or women. But she argues in Haslanger 2006 that, given metasemantic externalism, the counterintuitive nature of this consequence is unproblematic: just as 'water' may mean H_2O even if I think it denotes one of four basic elements, 'woman' may denote a social category that ought to be empty even if I think it denotes an unproblematic biological category.

On the other hand, if metasemantic externalism is right about the role played by *external* factors in determining meaning, the conceptual engineer may have insufficient control over meaning to make her project a practical possibility.[1] For example, Herman Cappelen concludes from metasemantic externalism that

> it is an illusion to think that any individual or group has any significant degree of control of the reference-fixing facts. If we are not in control of the reference-fixing facts, then we're not in control of conceptual engineering because it requires us to change the reference-fixing facts. (2018, p. 74)

If Cappelen is right then, far from justifying conceptual revisions, metasemantic externalism puts conceptual engineering out of reach.

It is important to get clear on the relationship. Metasemantic externalism is, amongst contemporary philosophers, perhaps the dominant view of how linguistic expressions obtain their meanings; and conceptual engineering is receiving increasing attention as a potentially important philosophical method. A supportive relationship benefits both; an antagonistic one benefits neither.

In the present paper, however, I argue for a relationship far looser than envisaged by either Haslanger or Cappelen: metasemantic externalism, I suggest, may bear very little on conceptual engineering. While metasemantic externalism principally operates at the level of *semantic*-meaning, conceptual engineering can largely operate at the level of *speaker*-meaning. In general, metasemantic

---

[1] Koch (2018) calls this the *externalist challenge*.

externalism neither justifies conceptual revisions, nor puts conceptual engineering out of reach. I call this the *Speaker-Meaning Picture of Conceptual Engineering*.[2]

The Speaker-Meaning Picture has two important benefits. Firstly, it makes conceptual engineering often *easy* in practice, as we have significant control over speaker-meaning. Secondly, the picture suggests a new, normative response to the well-known objection that conceptual engineering serves only to change the subject: if a conceptual revision yields a good thing for a speaker to mean, the conceptual engineer may thereby be *permitted* to change the subject. On the Speaker-Meaning Picture, changing the subject is not always a bad thing to do.

I begin (§§2–3) by critically examining Haslanger's (2006) and Cappelen's (2018) views on the relationship between conceptual engineering and metasemantic externalism. I then develop the Speaker-Meaning Picture (§4) and close by considering 'change of subject' objections (§5).

Before proceeding, three comments are in order. Firstly, conceptual engineers are rarely explicit about what, in the relevant sense, 'concepts' are. As Cappelen notes, '[t]hey often just talk about "concepts" […] and then leave it at that' (2018, p. 141). Nonetheless, there are two assumptions about 'concepts' that are typically (implicitly) made: that they can be delineated via a characterisation of their content, perhaps in the form of a definition; and that they are closely related to, or can be treated as, meanings of linguistic expressions.[3] Thus, one might *revise* WOMAN in the sense that one: introduces a revisionary definition of the form 'S *is a woman* iff …'; and subsequently uses the term 'woman' in accordance with that definition.

---

[2] Various recent views of conceptual engineering can be interpreted along these lines. For example, Brun (2016, p. 1217) takes concepts to be word/rules-of-use pairs, Prinzing (2018, p. 858) takes them to be cognitive tools, and Kitsik (2018, §3) takes concepts to be ways of using expressions in thought or talk. The Speaker-Meaning Picture makes explicit and develops what I take to be the core idea underlying such views.

[3] In recent work, Haslanger (forthcoming) understands concepts to be capacities to partition logical space. Both assumptions hold of this view. (While she takes a concept's content—a partition—to be inessential to it, she delineates concepts by characterising their contents.) In line with my comments below, I suggest that nothing would be lost if Haslanger recast the account as concerning (some variety of) *meaning*.

Secondly, in light of the above, there might not be a substantive role in conceptual engineering for anything we should call 'concepts'. Conceptual engineering may ultimately involve *changing the meanings* of linguistic expressions, or *changing how we use* those expressions, or *changing the definitions we associate with* those expressions, or something similar. From such a perspective, talk of 'concepts' in conceptual engineering is nothing more than a *façon de parler*, and 'conceptual engineering' is something of a misnomer.[4] Henceforth, I avoid talk of 'concepts' where appropriate.

Thirdly, I discuss metasemantic externalism throughout. *Metasemantics* is the body of theory concerning how linguistic expressions obtain their meanings. This is to be contrasted with *semantics*, the body of theory concerning what linguistic expressions mean. A metasemantic theory is *externalist* if it states that factors external to the individual speaker play a role in determining what linguistic expressions mean. There are a wide variety of externalist metasemantic theories, discussion of which goes far beyond the scope of this paper. Herein I follow Haslanger and Cappelen in focusing on theories in the tradition of Putnam (1975), Burge (1979) and Kripke (1980). Specific variants of metasemantic externalism will be introduced in due course.

2. Haslanger's Ameliorative Project

Haslanger aims 'to develop accounts of gender and race that will be effective tools in the fight against injustice' (2000, p. 36). To this end, she seeks to revise gender and race concepts. Herein I focus solely on gender concepts, although what I say applies *mutatis mutandis* to race concepts.

In Haslanger's 2006 formulation of the project, she appeals to metasemantic externalism to justify the proposed revisions. She begins by distinguishing three concepts that may be associated with any given term:

---

[4] Cf. Cappelen 2018, p. 3–4.

(a) the concept we take ourselves to by employing, (b) the concept that best captures the type we are concerned with, and (c) the type we ought [all-things-considered] to be concerned with […]. (p. 96)

Haslanger calls (a) the *manifest* concept, (b) the *operative* concept, and (c) the *target* concept.[5]

Here is a simple example. Suppose I think of myself as a beer connoisseur, and I take myself to apply 'tasty beer' to whatever beer has a rich and complex taste. However, not having a refined palate, I in fact apply 'tasty beer' to whatever beer has an overwhelmingly hoppy taste. Finally, suppose that, given how 'tasty' is used more generally, I ought all-things-considered to apply 'tasty beer' to whatever beer I in fact prefer the taste of—which, we might imagine, is neither rich and complex nor overwhelmingly hoppy. In this example, when using 'tasty beer': I take myself to be applying the concept *rich and complex*, so that is the manifest concept; the type I am in fact concerned with is *overwhelmingly hoppy*, so that is the operative concept; and I ought rather to be concerned with *my gustatory preference*, so that is the target concept.

According to Haslanger, '[in] the best cases, all three (my manifest, operative, and target concepts) will coincide' (p. 99). So when they come apart, something is to be done. For example, we might replace the manifest concept with the operative concept (a 'descriptive project'), or replace both the manifest and operative concepts with the target concept (an 'ameliorative project'), or eliminate the term altogether (an 'eliminativist project'). In the case of gender terms, Haslanger favours amelioration: we should change how language users understand themselves to be using those terms, as well as their actual practice, so that both self-understanding and practice conform to the target concept.

This leads to seemingly counterintuitive consequences. Haslanger believes that, all-things-considered, we ought to define gender concepts in terms of the social privilege and subordination that different groups experience (2000, p. 42). Roughly: a *man* is someone who is privileged on account of their perceived male bodily features; and a *woman* is someone who is subordinated on account of their perceived female bodily features. However, sex-marked privilege and subordination should be

---

[5] The 'all-things-considered' clause appears in Haslanger's definition of 'target concept', stated on p. 99.

eliminated so, given these revisionary definitions, we should aim 'to bring about a day when there are

no more women' (p. 46)—*mutatis mutandis* for men. This surprising conclusion leads to a well-

known worry: hasn't Haslanger simply changed the subject?[6]

To justify her conceptual revisions, defending them against this worry, Haslanger develops

the following variant of metasemantic externalism:

> *Objective Type Externalism*: Terms/concepts pick out an objective type, whether or not we can
>
> state conditions for membership in the type, by virtue of the fact that their meaning is
>
> determined by ostension of paradigms (or other means of reference-fixing) together with an
>
> implicit extension to things of the same type as the paradigms. (2006, p. 110)

Here, *natural kinds* (such as H_2O) are one kind of objective type; but, equally, *social kinds* (such as

those subordinated on account of perceived female bodily features) are another kind of objective type.

Thus, according to Haslanger, Objective Type Externalism helps us to understand the ameliorative

project.

> In such cases, we have perhaps a partial or vague understanding of the manifest concept, and
>
> the operative concept picks out a relatively heterogeneous set, but nonetheless we can say that
>
> there is something we mean, an objective type we are approaching. As before, I will use the
>
> term 'target concept' for the concept that is plausibly what we are getting at, even if we poorly
>
> understand it; the target concept is the object of ameliorative analysis. (p. 109).

Haslanger concludes that:

> [s]ocial constructionists can rely on externalist accounts of meaning to argue that their
>
> disclosure of [a] target concept is not *changing the subject*, but better reveals what we mean.
>
> (p. 110, emphasis in original).

---

[6] See Haslanger 2006, pp. 103ff. I discuss this objection in §5.

That is, according to Haslanger, Objective Type Externalism justifies the counterintuitive consequences of the ameliorative project.[7]

Contrary to Haslanger's claims, however, Objective Type Externalism does not 'reveal' target concepts. If we apply Objective Type Externalism to some term X, it will 'reveal' whichever concept picks out the objective type of whatever we paradigmatically describe as X. But, in general, that is *not* the *target* concept, the concept we *ought all-things-considered* to be using, but rather the *operative* concept, the concept that *best captures* the type we are concerned with.[8] The upshot is this: Objective Type Externalism can be used to justify an ameliorative project *only when the target concept and operative concept coincide*.

For gender terms, do target and operative concepts coincide? A full answer would require empirical data on who we describe as 'man' and 'woman', and a thorough normative analysis of how we ought all-things-considered to use 'man' and 'woman', neither of which can be provided here. However, it is helpful to consider an adaptation of the twin-earth thought experiment.

> There is a distant planet, very similar to earth, which we may call *twin-earth*. On twin earth, there are plants, animals and humanoids that are biologically identical to those on earth. However, the twin-humans are socially different to humans, just in the sense that: those perceived as having male bodily features are not privileged (or subordinated) on that basis; and those perceived as having female bodily features are not subordinated (or privileged) on that basis.

---

[7] Cappelen argues that 'amelioration as revelation undermines the basic idea behind ameliorative projects […]. What we thought was amelioration (i.e., improvement of our representational devices) turns out to be revelation of something that was there all along' (2018, p. 80). However, the fact that Haslanger *engineered* the concept that she takes Objective Type Externalism to 'reveal' suggests that, *pace* Cappelen, Haslanger can in principle build 'revelation' into a conceptual engineering project. Herein, I leave this objection aside.

[8] On p. 109, Haslanger writes: 'As before, I will use the term 'target concept' for the concept that is plausibly what we are getting at, even if we poorly understand it'. *Pace* Haslanger, this is *not* how she used the term before.

Intuitively, one might describe twin-earth by saying 'there is equality between men and women on twin-earth', or 'women have things better on twin-earth than on earth'. Very few would intuitively describe twin-earth by saying 'there are no men or women on twin-earth'. This suggests that our use of 'man' and 'woman' does *not* coincide with Haslanger's definitions—that the operative concept associated with those terms does not coincide with what Haslanger takes to be the target concept.

Twin-earth thought experiments are not decisive, and appeals to intuitions must be undertaken with care. But, nonetheless, the above gives us an initial reason to doubt that the operative concept associated with gender terms coincides with Haslanger's definitions. And, given the fact that Haslanger's definitions are explicitly revisionary, such a conclusion ought not to be surprising. Absent the kind of empirical and normative details that would be required to show otherwise, we should doubt that Haslanger's revisionary definition coincides with the operative concept.

Given Haslanger's characterisation of the target concepts associated with 'man' and 'woman', then, we should doubt that Objective Type Externalism 'reveals' those target concepts as the meanings of those terms. As things stand, Haslanger cannot use Objective Type Externalism to justify her ameliorative project.

3. Cappelen's Austerity Framework

According to Cappelen's *Austerity Framework* (2018), conceptual engineering doesn't involve concepts, and barely involves engineering. First, driven by concerns about the nature of concepts combined with an appeal to simplicity (pp. 141ff), Cappelen writes that 'conceptual engineering should be seen as having as its goal to change extensions and intensions' (p. 61)—or, more precisely, 'changes in extensions that are driven by changes in intensions' (p. 62).

To understand how intensions of expressions change, he appeals to a metasemantic theory. Cappelen appeals to a variant of metasemantic externalism with multiple components. Here are the key aspects of his view.

Cappelen's Externalism:

- External factors—the environment, experts, historical chains of use, etc.—play a principal role in determining the intensions and extensions of linguistic expressions.

- There is a genuine possibility that speakers are in massive error about those intensions and extensions.

- 'In most cases the detailed mechanisms that underpin particular instances of conceptual engineering are too complex, messy, non-systematic, amorphous, and unstable for us to fully grasp or understand.' (p. 72)

- 'The process of conceptual engineering is governed by factors that are not within our *control*: no individual or group has a significant degree of control over how meaning change happens.' (*ibid*., emphasis in original)

On this view, conceptual engineering is *incredibly* difficult. We do not know in detail *how* meaning-change occurs and, even if we did, we would not be in control of the relevant mechanisms. If one seeks to make a specific change to the intension of an expression, one is incredibly *unlikely* to succeed. Given Cappelen's Externalism, conceptual engineering is, in practice, effectively impossible.[9]

Let us accept Cappelen's Externalism for the sake of argument.[10] I argue that: in certain cases, we should not understand conceptual engineering to involve changes of intension; and in such cases, Cappelen's Externalism notwithstanding, conceptual engineering can be *easy* in practice.

The cases in question are those in which someone makes explicit that they will use a term nonstandardly within a local context. For example:

---

[9] Cappelen is *not* engaged in a debunking project. He remains in favour of conceptual engineering—just as he is in favour of trying to reduce poverty in Bangladesh despite it being, for most of us, effectively impossible in practice (2018, pp. 74–75, 200–201). Nevertheless, the practical impossibility of conceptual engineering is a very unwelcome consequence for the conceptual engineer.

[10] As Koch (2018) argues, Cappelen's Externalism is probably too strong: desired changes in meaning can be brought about by groups of speakers who co-ordinate their language use over time.

The Social Media Argument. There are two types of people who do not use social media. There is the person who is opposed to social media, often having had (and closed) a *Facebook* account due to its negative effects on society. Call her the *Luddite*. In contrast, there is the person who has no specific objection to social media, but simply prefers to interact using more traditional means. Call her the *Traditionalist*. To fight against the spread of social media, we should not seek to increase the number of Luddites: the Luddite is reacting to the popularity of social media, and reactionaries stand little chance of changing mass movements. Rather, we should seek to increase the number of Traditionalists; we should make a movement of traditional means of social interaction.

There are two immediate points to make about the Social Media Argument. Firstly, it involves an *attempt* at conceptual engineering. In the opening sentences, nonstandard definitions are given for 'Luddite' (viz. *one who does not use social media and is opposed to it due to its negative effects*) and 'Traditionalist' (viz. *one who does not use social media due to a preference for traditional means of interaction*). It is clear in context that these definitions have been constructed for a specific purpose (viz. formulating the Social Media Argument). This is an archetypal attempt at conceptual engineering.

Secondly, whatever else might be said of the argument, it is *understandable*: the reader can grasp the author's argument. And the argument is understandable *because* the author successfully conveys to the reader how to locally interpret 'Luddite' and 'Traditionalist'. If those terms were interpreted according to standard usage, the argument would be almost incomprehensible: there was no social media at the time of the Luddites and, if 'Traditionalist' has a stable meaning at all, it does not distinctively connote social media.

Now, let us focus on a scenario in which a reader straightforwardly grasps the author's argument, just as the author intended. How should we describe such a scenario? Cappelen (2018, pp. 75–77) considers similar cases, offering two possibilities.[11] Applied to the present scenario:

---

[11] Cappelen makes brief comments about speaker-meaning in this section, which I address in §4.1.

10

Meaning-Change. Conceptual engineering involves changing intensions. The author of the Social Media Argument locally changes the intension of 'Luddite' and 'Traditionalist'. The reader interprets those terms correctly. Thus, successful conceptual engineering leads to successful communication.

Folie-à-Deux. Conceptual engineering involves changing intensions. The author of the Social Media Argument attempts, and fails, to locally change the intension of 'Luddite' and 'Traditionalist'. The author believes herself to have been successful, and subsequently misuses 'Luddite' and 'Traditionalist' on that belief. The reader misinterprets 'Luddite' and 'Traditionalist' accordingly. Thus, failed conceptual engineering and complementary errors lead to successful communication.

Now, Cappelen's Externalism rules out Meaning-Change: given Cappelen's Externalism, we lack sufficient control over meaning to make local changes to the intensions of 'Luddite' and 'Traditionalist'. Thus, in the present scenario, Cappelen would endorse Folie-à-Deux.

However, this is a false dichotomy. Rather than accepting Folie-à-Deux, we might deny that, here, conceptual engineering involves changing intensions. Note first that Cappelen would *accept* that the author (mistakenly) uses 'Luddite' and 'Traditionalist' *as if* they have nonstandard definitions (p. 76). But Cappelen would maintain that *merely* 'using-as-if' in this way amounts to a *failed attempt* at conceptual engineering (*ibid.*). *Pace* Cappelen, I suggest we deny the latter claim and offer a more charitable description of the scenario:

Use-As-If. It is sufficient for conceptual engineering that a term be explicitly used *as if* it has a newly constructed, nonstandard definition. The author makes explicit that she will locally use 'Luddite' and 'Traditionalist' *as if* they have such definitions, and uses those terms accordingly. The reader recognises the author's indication, and locally interprets 'Luddite' and 'Traditionalist' *as if* they have those definitions. Thus, this instance of conceptual engineering is successful, leading to successful communication.

11

Note that, here, 'using X *as if* it has definition *d*' is intended to be neutral over whether X *in fact* has definition *d*: Use-As-If does not rule out the possibility that the author succeeds in locally changing the meaning of 'Luddite' and 'Traditionalist'.[12]

Given Cappelen's Externalism, Use-As-If is preferable to Folie-à-Deux for at least three reasons. Firstly, Folie-à-Deux is uncharitable: it describes the author as attempting the practically impossible, despite a more reasonable interpretation in the vicinity. There is no clear reason to insist that the author intends *to locally change the intensions of 'Luddite' and 'Traditionalist' to something nonstandard*, rather than *to explicitly use 'Luddite' and 'Traditionalist' as if they have nonstandard definitions*.

Secondly, Folie-à-Deux is overcomplicated: it interprets a simple case of communication by imputing complementary errors to author and reader alike. This forces us to distinguish between what both parties *think* they are doing and what they are *actually* doing, an additional layer of description. However, as captured by Use-As-If, this scenario in fact involves a run-of-the-mill instance of communication, without belief being divorced from action in this way.

Thirdly, Folie-à-Deux obscures why communication is successful in this scenario. In general, successful communication involves an audience grasping whatever a speaker intends to convey. Given Folie-à-Deux, the reader grasps what the author intends to convey *despite* the author's efforts: the author makes a mistake and is to some extent *lucky* that the reader makes a complementary error. But this obscures the obvious: in the present scenario, communication is in fact successful *because* the author is explicit about *how she will use* 'Luddite' and 'Traditionalist'. This is captured straightforwardly by Use-As-If.

Given Use-As-If, conceptual engineering can be *easy* in practice. It is within a careful author's control to clearly construct and articulate a nonstandard definition for a term and use that term accordingly; and, once articulated so, it is within the careful reader's control to locally interpret the relevant term as if it has that definition. So, even given Cappelen's Externalism, conceptual engineering is *not*, in practice, effectively impossible.

---

[12] Thus Use-As-If is compatible with, e.g., Ludlow (2014) and Plunkett and Sundell (2013).

I have argued that, in certain cases, we should understand conceptual engineering *not* to involve changing extensions and intensions. This is a *normative*, *conceptual* claim about *what we should count as* conceptual engineering. Moreover, given Cappelen's line of argument in *Fixing Language*, he might well agree. Cappelen does not argue *for* the claim that conceptual engineering involves changing extensions and intensions. Rather, he argues *against* seeing conceptual engineering as aiming to engineer *concepts* (2018, pp. 141ff), and offers explanations of why conceptual engineers *think* they are engineering concepts (pp. 85–93). His appeal to extensions and intensions is then justified by simplicity:

> I appeal to non-controversial ingredients that you already have at your disposal, I leave out all
> the controversial machinery, and I can explain the same phenomena. Simplicity recommends
> the Austerity Framework […]. (p. 142)

However, the present suggestion appeals only to non-controversial ingredients you already have at your disposal, leaves out all the controversial machinery, and captures such cases as the Social Media Argument *better than* the Austerity Framework. By Cappelen's own lights, then, there are grounds for accepting that conceptual engineering does not always involve trying to change extensions and intensions—that metasemantic externalism does not make conceptual engineering effectively impossible in practice.

4. The Speaker-Meaning Picture of Conceptual Engineering

In the Social Media Argument, conceptual engineering plausibly consists in explicitly using terms *as if* they have newly constructed definitions. I now develop the idea in terms of *speaker-meaning*. The result, I suggest, is an attractive approach to conceptual engineering more generally.

4.1 Speaker-Meaning

The distinction between speaker-meaning and semantic-meaning is familiar.[13] *Speaker-meaning* is

what a *speaker* means when she utters some words. It is closely tied to a speaker's intentions: in

ordinary conversational exchanges, what a speaker means when she utters some words is just what

she intends to convey by her utterance.[14] In contrast, *semantic-meaning* is what the uttered *words*

mean. For concreteness, we can think of semantic-meaning as closely tied to the idea of a linguistic

community: what a word means is governed in some way by the linguistic community to which the

speaker belongs. This situation is complicated by the fact that even monolingual speakers might

belong to many linguistic communities: the linguistic community of anglophone philosophers, for

example, might be distinct from that of the English village in which I grew up. But, insofar as we can

understand a speaker to be speaking *qua* member of a particular linguistic community, I will assume

the idea is workable.

Nothing that follows depends on a specific theory of the nature of speaker-meaning or

semantic-meaning. For concreteness, I assume that both can be modelled as intensions and/or

extensions. I assume that intensions and extensions can be characterised using italics (e.g. 'grass is

green' semantically-means *grass is green*; I speaker-mean *grass* by 'grass'), and treated as objects

(e.g. 'snow is white' semantically-means m; I speaker-mean m' by 'snow'). The model is an

idealisation: for most terms, at least, I doubt that our intentions suffice to determine specific

intensions and extensions. Nonetheless, it is sufficient for present purposes.

Speaker-meaning and semantic-meaning can come apart. Suppose that, upon opening the

curtains to another grey, miserable day, Sophie utters 'it's another *beautiful* day' with faux

enthusiasm. Plausibly, Sophie intends to convey *that it's another miserable day*, so that is what she

---

[13] See e.g. Grice 1989.

[14] We cannot define *speaker-meaning* in this way generally. For one example, there are cases in which speakers *mean*

something without an audience (and thus without an intention to convey anything at all). See Grice 1989 for attempts to deal

with this and related issues.

speaker-means. However, supposing that Sophie belongs to a regular community of English speakers, 'it's another beautiful day' semantically-means *that it's another beautiful day*.

How *far* they can come apart is plausibly limited by our capacity to intend. In general: to speaker-mean *m′* by X, one must *intend* to convey *m′* by X; and to *intend* something, one must take it to be achievable. For example, Donnellan writes:

> In the analysis of meaning given by Grice, a speaker means something by an utterance when he has a certain complex kind of intention […]. And what the speaker means is determined by the content of that intention. Whether he can form that intention, however, may depend upon what expectations he has about his audience and their ability to grasp his intention. (1968, p. 212)

So, in an ordinary conversational exchange, Sophie cannot speaker-mean *that Paris is in France* by 'it's another *beautiful* day' because she cannot form the requisite intention; and she cannot form the requisite intention because, it being an ordinary conversational exchange, she does not expect her audience to interpret her that way. On this line of thought, speaker-meaning can come apart from semantic-meaning to the extent that a speaker can expect her audience to interpret her accordingly.

But now, as Donnellan (p. 213) goes on to suggest: if we are *explicit* about how we plan to use a term, then, even in quite extreme cases, we can reasonably expect our audience to interpret us accordingly. Consider the child who tells her parents that she will swap 'yes' and 'no'. For a limited time, she can then expect to convey, and can thus speaker-mean, *no* by 'yes' and *yes* by 'no'. Likewise for the author who explicitly offers a nonstandard definition for 'Luddite': within local contexts, she can subsequently expect to convey, and can thus speaker-mean, something nonstandard by that term. In both cases, the audience is likely to understand, and communication is likely to be successful.

It does not follow that we can speaker-mean anything we like by any term whatsoever: our expectations of the audience may be affected by the audience's linguistic sophistication, the complexity of the definition, the linguistic context, the political and social climate, and so on. But if the speaker is explicit about her intentions then, in an appropriate context, with an appropriate

audience, speaker-meaning can deviate from semantic-meaning in new and substantial ways, and the audience will be able to grasp that speaker-meaning.

Underlying these comments is the view that successful communication relies principally on conveying speaker-meaning. One might ask whether semantic-meaning thus drops out of the picture altogether. Consider Davidson's discussion in *A Nice Derangement of Epitaphs* (2005). On Davidson's view, we come to every conversation armed with an *audience-relative* set of expectations as to how our words will be interpreted (encoded in the *prior* theory), revising our expectations throughout the conversation (yielding sequential *passing* theories). Davidson claims that '[what] must be shared for communication to succeed is the passing theory' (p. 102). And, due to differences in how we use language and prejudge others to do so, he claims that 'most of the time prior theories will not be shared' (p. 103). Davidson concludes that neither passing theories nor prior theories are conventional, and thus that there are no conventional semantic-meanings (or *first meanings*).

Now, in Davidson's terminology, I have effectively suggested: that a speaker can effect a co-ordinated shift in conversational participants' passing theories, with respect to a given term, by being explicit about her intentions; and that, as Davidson claims, this is sufficient for successful communication. But this is *neutral* on whether *prior* theories are shared or conventional. If Davidson is *right* that prior theories are not shared, then conventional first meanings do indeed drop out of the picture. It might then be harder for speakers to explicitly convey their intentions successfully—especially to wider audiences—but only insofar as this would be true of communication more generally. In contrast, if, *pace* Davidson, prior theories are conventional, then conventional first meanings remain *in* the picture. Either way, a speaker who carefully makes her intentions explicit can successfully communicate via novel speaker-meanings. Regardless, given my broader concerns, I continue to assume herein that there *are* semantic-meanings.

Throughout this subsection, I have argued that we have control over speaker-meaning. Cappelen disagrees:

> The content of what's called 'speaker's meaning' [i.e. what the speaker has in mind] is just as externalistically determined as linguistic meaning: we have no more control over that content than we have over what we say when we utter sentences in public language. (2018, p. 76)

The thought, insofar as I understand it, is that mental content is externalistically determined, in broadly the same way as semantic-content; and so we have no more control over the content of our intentions than over semantic-content. But, even if I lack control over the semantic-meanings of sentences, I can control which sentences I utter. I might be unable to make 'it's hot' semantically-mean *that it's cold*, but I can utter 'it's cold' instead. Likewise, even if mental content is externalistically determined, I have significant control over what I intend. For example, when uttering 'Luddite', I control whether I intend to convey *Luddite* or, alternatively, *one who does not use social media and is opposed to it due to its negative effects*. *Pace* Cappelen, we have significant control over what we intend to convey and, *a fortiori*, over speaker-meaning.

4.2 The Picture

I propose the following:

> <u>The Speaker-Meaning Picture of Conceptual Engineering</u>
>
> *Descriptive claim*. Suppose that a term, X, semantically-means m.[15] Then, *S* can engage in conceptual engineering with respect to X by
>
> - making explicit, perhaps by means of a definition, that she will speaker-mean m' by X in local contexts,
> - speaker-meaning m' by X within local contexts, and
> - (implicitly) presenting m' as a good thing to speaker-mean by X within local and relevantly similar contexts,
>
> for some m' ($\neq m$) that *S* has constructed/engineered for this purpose, where *relevant similarity* is fixed by how *S* (implicitly) presents her goals.

---

[15] I allow m to be *null*, representing the fact that X may have no (semantic-)meaning at all.

Normative claim. This is a good approach to conceptual engineering.

The descriptive claim delineates what I call the *speaker-meaning approach* to conceptual engineering. To illustrate, recall the Social Media Argument. The author of the argument begins by making explicit what, in formulating her argument, she will speaker-mean by 'Luddite' and 'Traditionalist'. She then uses those terms accordingly. In doing so, she implicitly presents her definitions as good things to speaker-mean by those terms in contexts in which one is formulating that argument. Careful readers will understand the author's intentions easily and interpret her accordingly. And those who decide to relay or repudiate the argument are likely to adopt her usage when formulating the argument for that purpose, explicitly deferring to the author's definitions in those contexts. For her purposes, then, the author has plausibly engaged successfully in conceptual engineering.

This instance of conceptual engineering is narrow in scope. The new definitions are presented as good things to speaker-mean by 'Luddite' and 'Traditionalist' in a very limited range of contexts. In contrast, Haslanger is interested in how gender terms are used *in general*. Such *wide-scope* conceptual engineering can be carried out within the speaker-meaning approach in the same way as *narrow-scope* conceptual engineering.

For example, regardless of what exactly Haslanger (2000) had in mind when first introducing her project: she explicitly introduced revisionary definitions for 'man' and 'woman'; she subsequently used those terms accordingly within that article;[16] and she presented those definitions as good things to speaker-mean by those terms in general—*good*, because making sex-marked privilege and subordination explicit in our gender-talk may yield a more gender-equal society. On the present picture, she was thus engaged in conceptual engineering—regardless of whether she *also* intended to revise semantic-meanings or replace concepts.

So construed, however, Haslanger has perhaps not yet, given her goals, been successful. I concede that Haslanger's readers were largely able to interpret her uses of 'man' and 'woman'; and some commentators have perhaps deferred to Haslanger's definitions, as when Saul writes that '[it] is

---

[16] An arbitrary example: 'On my analysis women are those who occupy a particular *kind* of social position' (2000, p. 45, emphasis in original).

built into Haslanger's definitions […] that women are systematically subordinated to men' (2006, p. 122). But Haslanger aimed to influence *people in general*, and this has *not* been the case: most people do *not* defer to Haslanger's definitions when using 'man' and 'woman'. Of course, it is incredibly difficult to shift people's intentions in this way. But such wide-scope projects will be difficult on any plausible approach to conceptual engineering.

Regardless, the speaker-meaning approach is a *good* approach to conceptual engineering. Firstly, in most cases, it gives conceptual engineers what they want. Carefully articulated speaker-meanings can convey new information and new ideas. They can be highly specific, suitable for drawing important and subtle distinctions not encoded in terms' semantic-meanings. They can give people new ways to think about, or 'carve up', certain phenomena. And perhaps, if enough people change what they ordinarily speaker-mean by the terms, they may shift prevailing ideas about those phenomena.

Moreover, success on the speaker-meaning approach is often easy. Situations commonly arise in which *local* conceptual engineering is called for. Perhaps the topic under discussion requires unusual precision, new distinctions, or something else. Such situations are ubiquitous in academic papers, medical texts, legal documents, etc., and arise more generally in conversations about subtle and demanding topics. When they arise, the careful author (or speaker) articulates how she intends to use relevant terms, and careful audiences interpret her accordingly. Of course, success is not guaranteed—definitions can be faulty, people can misuse their definitions, and so on. But, with sufficient care, the conceptual engineer will typically be successful in narrow-scope cases.

Even in wide-scope cases, the speaker-meaning approach admits of a clear, general strategy for achieving success. In these cases, the conceptual engineer aims to change a specific aspect of linguistic behaviour far more widely. So she should publicise the new definition, building public momentum for its adoption. At the outset, this may involve seeking to convince a high-profile group of users to defer to her definition. Then, over time, she might broaden that group, selectively targeting individuals who are influential in different communities. She will need simple and intuitive

definitions, compelling arguments and powerful examples, as well as time, influence and luck.[17] But we can recognise how, in the right circumstances, this could be achieved.

Finally, on the speaker-meaning approach, the conceptual engineer obtains these benefits without commitment to a substantive metasemantic theory. This contrasts sharply with approaches to conceptual engineering that seek to change *semantic*-meanings. To see this, consider a view developed by Peter Ludlow (2014). According to Ludlow, everyday communication relies on building *microlanguages* by collaboratively *modulating* and *sharpening* semantic-meanings. Much of this happens implicitly. For example, he recalls playing with an action figure, before the term 'action figure' had been coined (p. 9). His father asked 'why are you playing with that doll?', to which Ludlow replied 'it's not a doll'. The dispute, Ludlow claims, was about what counted as a doll: they were implicitly negotiating the semantic-meaning of 'doll' within their microlanguage.

Against Ludlow's theoretical backdrop, conceptual engineering is continuous with ordinary conversational practices. We can merely think of the conceptual engineer as modulating or sharpening a semantic-meaning *explicitly*. In practice, conceptual engineering on this approach is very similar to conceptual engineering on the speaker-meaning approach: in narrow-scope cases, the conceptual engineer articulates, and 'litigates' for, her proposed modulation or sharpening; in wide-scope cases, she aims to convince people to accept the proposal more generally.

The difference is in the theory. On the latter view, the conceptual engineer aims to modulate or sharpen the *semantic*-meaning of a term—where this is possible because semantic-meanings are *underdetermined*, terms do not have *privileged* or *primary* semantic-meanings, and *ordinary conversational participants* have *control* over what their words semantically-mean (pp. 80–84).[18] This

---

[17] All of this is compatible with the conceptual engineer's proposed definition being defective—encoding morally or socially problematic tropes, or failing to cut nature at its joints, or something else. As an anonymous referee puts it, the speaker-meaning approach could lead to 'conceptual deterioration'. In pushing for new definitions, then, the wide-scope conceptual engineer takes on significant moral, social and epistemic responsibility.

[18] A similar view has been developed by Plunkett and Sundell (2013). They do not consistently distinguish between semantic-meaning and speaker-meaning, and so the mechanisms underlying their view are not entirely clear. As they

approach to conceptual engineering thus presupposes at least three substantive—and controversial—metasemantic theses.[19] In contrast, a speaker can intend to convey something new by a term, she can make this intention explicit, and she can subsequently use that term and be interpreted accordingly, *regardless* of how—and by what—semantic-meanings are determined.

Despite being similar in practice, then, these approaches to conceptual engineering have very different theoretical commitments. The point generalises. On the Speaker-Meaning Picture, conceptual engineering is often easy, admits of a general strategy in difficult cases, and operates independently of metasemantics. If we think of conceptual engineering as changing *semantic-*meanings, however, the former benefits are available only to those taking on substantive metasemantic commitments.

I now turn to an additional benefit of the Speaker-Meaning Picture: it suggests a new, normative response to the well-known objection that conceptual engineering serves only to change the subject.

5. 'Change of Subject' Objections

Perhaps the best-known objection to conceptual engineering is that, at least in philosophy, it serves only to change the subject.[20] The problem is typically attributed to Strawson who, in discussing Carnap's (1950) method of explication, wrote that

> […] typical philosophical problems about the concepts used in non-scientific discourse cannot
> be solved by laying down the rules of use of exact and fruitful concepts in science. To do this
> last is not to solve the typical philosophical problem, but to change the subject. (1963, p. 506)

---

officially present the view, it presupposes that ordinary conversation is sufficient to locally change semantic-meanings in a controlled way (pp. 8–9, 26–28). So, as officially presented, the following comments apply *mutatis mutandis* to their view.

[19] Cf. Cappelen 2018, pp. 163–179.

[20] The objection has been discussed at length in the recent literature. See e.g. Cappelen 2018; Haslanger 2006; Pinder 2017; Prinzing 2018; Schupbach 2017.

Strawson's focus is on the effects of replacing ordinary 'concepts' with those designed to be useful in science. But, as we saw in §2, the point generalises: if Haslanger addresses inequality between men and women by revising the definitions of 'man' and 'woman', she may seem only to succeed in changing the subject.

At face value, changing the subject in philosophy can be problematic in at least two kinds of case. Firstly, suppose that S is tackling a philosophical problem concerning x. If S changes the subject to y ($\neq x$) in that context, she equivocates and fails to engage with the problem. For example, in epistemology, Gettier cases pose problems for certain analyses of *knowledge*. So, at face value, one cannot tackle those problems by speaker-meaning something other than *knowledge* by 'knowledge'.[21] Secondly, suppose that x is something we care about, in the sense that we either value or disvalue x, and that S is examining x. If S changes the subject to y ($\neq x$) in that context, she fails to engage with the very thing should make her discussion worthwhile. For example, we care about *what is good for us*. So, at face value, one who examines *what is good for us* by speaker-meaning something nonstandard by '(nonmoral) good' simply misses the point.[22] In such cases, subject-preservation appears to be a good thing—and conceptual engineering appears problematic.

One possible response is to provide a criterion for subject preservation. Jonah Schupbach (2017) argues that, when replacing one concept with another, one preserves the subject if, in the contexts at hand, the two concepts are extensionally very similar. As we can model both speaker-meaning and semantic-meaning as having extensions, Schupbach's response, if successful, carries over to the Speaker-Meaning Picture.

Another possible response is to emphasise cases in which we *want* to change the subject. For example, Joshua Glasgow (2009) proposes that we replace 'race-talk' with 'race*-talk'. According to Glasgow, race-talk purports to refer to something biological and is thus empty, whereas race*-talk refers successfully to wholly social objects. On Glasgow's view, *changing* the subject is a desideratum.

---

[21] Cf. Olsson 2015.

[22] See Railton 1989—though note that Railton argues at length that the proposal is '*tolerably* revisionist' (pp. 168–172).

The Speaker-Meaning Picture, however, is suggestive of an alternative response. Recall that, on the Speaker-Meaning Picture, the conceptual engineer presents *m′* as a *good thing to speaker-mean* by X, in local and relevantly similar contexts. Consider the following question: what makes *m′* a good or bad thing to speaker-mean by X?[23] Many factors may be relevant. Perhaps speaker-meaning *m′* would change the subject away from something we care about; perhaps it would lead to greater social equality; perhaps it would be a more precise way to talk about whatever X semantically-means; perhaps it would allow us to carve nature at its joints, or prevent us from doing so; perhaps it would facilitate, or hinder, giving explanations of key phenomena; and so on. Even when subject-preservation is desirable, it is just one, defeasible factor, which might in principle be outweighed by other factors.

Here, then, is a normative response to 'change of subject' objections. First, *S* is permitted to speaker-mean *m′* by X if *m′* is a good thing for *S* to speaker-mean by X in that context. Second, whether *m′* is a good thing for *S* to speaker-mean by X in that context depends on many factors, only one of which is whether *S* thereby changes the subject. So, even when subject-preservation is desirable, *S* may be permitted all-things-considered to changes the subject.[24]

Consider Kate Manne's (2017) ameliorative approach to misogyny. Manne motivates her discussion, in part, by asking 'Why is misogyny still a thing?' (p. xii) and 'What is misogyny?' (pp. 31–32, 41f). As these questions concern misogyny, and misogyny is something we care about, Manne has a *pro tanto* reason in this case to preserve the subject. However, Manne argues that the ordinary understanding of misogyny—as a general hatred of women—makes misogyny very difficult to identify (since it is difficult to establish that someone has a general hatred of women), and very rare (since few people have a general hatred of women). This is problematic:

---

[23] Simion (2017) discusses this issue in more detail, although not in the context of 'change of subject' objections.

[24] A similar response is considered by Cappelen (2018, pp. 122ff). According to Cappelen, some conceptual engineers might not care whether they change the subject, caring only about the *lexical effects* of the candidate semantic-meanings. There are two key differences. On the present suggestion, the *normative* question of whether the conceptual engineer is justified plays the central role, rather than what conceptual engineers care about; and it is not merely *lexical effects* that are considered, but *all things*.

Given that the notion of misogyny plausibly has a certain conceptual role to play of picking

out the most hostile and noxious facets of gender-based oppression, this would threaten to

deprive women of a suitable name for a potentially potent problem facing them. (p. 44)

Thus, Manne argues that we should *retain* the word 'misogyny', but use a *new* definition, designed to pick out the observable and common behaviours that constitute the most hostile and noxious facets of gender-based oppression (p. 63). Even supposing that this proposal would change the subject, Manne's reason for people to use 'misogyny' in this way, grounded in the demands of social and epistemic justice, plausibly outweighs the *pro tanto* reason to preserve the subject. If this is right, then Manne's definition is a good thing to speaker-mean by 'misogyny', and changing the subject is all-things-considered permitted.

Is this response available to other pictures of conceptual engineering? In principle, yes. First, notice that, in general, whether a speaker changes the subject depends on speaker-meaning not semantic-meaning. Suppose, for example, that 'the rabbit is out of the hat' semantically-means *the (relevant) rabbit is out of the (relevant) hat*; but that, on some occasion, Joan uses it to speaker-mean *the secret has been revealed*. On such an occasion, she counts as preserving the subject if the subject of conversation concerns a revealed secret, but not if the subject instead concerns rabbits. So subject-preservation tracks speaker-meaning; and whether *S* is permitted to change the subject on a given occasion is tied to what *S* is permitted to speaker-mean. And so, whether *S* is permitted to change the subject via a conceptual engineering project with respect to X depends on what she is permitted to speaker-mean by X. For the normative route outlined above to be available to a picture of conceptual engineering, then, that picture should imply that: whether one is justified in undertaking a conceptual engineering project with respect to X depends on what one is permitted to speaker-mean by X. Any picture meeting this constraint can be supplemented with this response to 'change of subject' objections.

6. Conclusion

The relationship between conceptual engineering and metasemantic externalism is, I have argued, much looser than envisaged by Haslanger and Cappelen. While metasemantic externalism principally concerns *semantic*-meaning, conceptual engineering can operate largely at the level of *speaker*-meaning. On this picture, conceptual engineering can be easy in practice; and it suggests a new, normative response to the objection that, in philosophy, conceptual engineering serves only to change the subject.[25]

References

Brun, Georg 2016, 'Explication as a Method of Conceptual Re-engineering', *Erkenntnis* 81(6), pp. 1211–1241.

Burge, Tyler 1979, 'Individualism and the Mental', *Midwest Studies in Philosophy* 4(1), pp. 73–122.

Cappelen, Herman 2018, *Fixing Language* (Oxford: Oxford University Press).

Carnap, Rudolf 1950, *Logical Foundations of Probability* (Chicago: University of Chicago Press).

Davidson, Donald 2005, 'A Nice Derangement of Epitaphs', in his *Truth, Language and History* (Oxford: Oxford University Press), pp. 89–107.

Donnellan, Keith 1968, 'Putting Humpty Dumpty Together Again', *Philosophical Review* 77(2), pp. 203–215.

Glasgow, Joshua 2009, *A Theory of Race* (Oxford: Routledge).

Grice, H.P. 1989, *Studies in the Way of Words* (Cambridge, MA: Harvard University Press).

Haslanger, Sally 2000, 'Gender and Race: (What) Are They? (What) Do We Want Them to Be?', *Noûs* 34(1), pp. 31–55.

—— 2006, 'What Good Are Our Intuitions? Philosophical Analysis and Social Kinds', *Proceedings of the Aristotelian Society Supplementary Volume* 80(1), pp. 89–118.

—— forthcoming, 'Going On, Not in the Same Way', in A. Burgess, H. Cappelen and D. Plunkett (eds.), *Conceptual Engineering and Conceptual Ethics* (Oxford: Oxford University Press).

Kitsik, Eve 2018, 'Explication as a Strategy for Revisionary Philosophy', *Synthese*, doi: 10.1007/s11229-018-1774-z.

Koch, Steffen 2018, 'The Externalist Challenge to Conceptual Engineering', *Synthese*, doi: 10.1007/s11229-018-02007-6.

Kripke, Saul 1980, *Naming and Necessity* (Cambridge, MA: Harvard University Press).

Ludlow, Peter 2014, *Living Words* (Oxford: Oxford University Press).

Olsson, Erik 2015, 'Gettier and the Method of Explication: A 60 Year Old Solution to a 50 Year Old Problem', *Philosophical Studies* 172(1), pp. 57–72.

Manne, Kate 2017, *Down Girl: The Logic of Misogyny* (Oxford: Oxford University Press).

Pinder, Mark 2017, 'On Strawson's Critique of Explication as a Method in Philosophy', *Synthese*, doi: 10.1007/s11229-017-1614-6.

Plunkett, David and Tim Sundell 2013, 'Disagreement and the Semantics of Normative and Evaluative Terms', *Philosophers' Imprint* 13(23), pp. 1–37.

Prinzing, Michael 2018, 'The Revisionist's Rubric: Conceptual Engineering and the Discontinuity Objection', *Inquiry* 61(8), pp. 854–880.

Putnam, Hilary 1975, 'The Meaning of "Meaning"', in K. Gunderson (ed.) *Language, Mind and Knowledge*, *Minnesota Studies in the Philosophy of Science (Vol. 7)* (Minneapolis: University of Minnesota Press), pp. 131–193.

Railton, Peter 1989, 'Naturalism and Prescriptivity', *Social Philosophy and Policy* 7(1), pp. 151–174.

Saul, Jennifer 2006, 'Gender and Race', *Proceedings of the Aristotelian Society Supplementary Volume* 80(1), pp. 119–143.

Schupbach, Jonah 2017, 'Experimental Explication', *Philosophy and Phenomenological Research* 94(3), pp. 672–710.

Simion, Mona 2018, 'The "Should" in Conceptual Engineering', *Inquiry* 61(8), pp. 914–928.

Strawson, P.F. 1963, 'Carnap's Views on Constructed Systems Versus Natural Languages in Analytic Philosophy', in P. Schilpp (ed.) *The Philosophy of Rudolf Carnap* (Lasalle, IL: Open Court), pp. 503–518.